

Fifth European Conference on Algorithmic Fairness (ECAF'26)

TIJL DE BIE, Ghent University, Belgium

MARYBETH DEFRANCE, Ghent University, Belgium

TOON CALDERS, University of Antwerp, Belgium

DENNIS NGUYEN, Utrecht University, Netherlands

The European Conference on Algorithmic Fairness (ECAF) aims to foster dialogue between researchers working on algorithmic fairness in the context of Europe's legal and societal framework, especially in light of the European Union's attempts to promote ethical AI and the turn to AI for the common good. ECAF welcomes submissions from multiple disciplines, including but not limited to computer science, law, philosophy, and social science, as well as interdisciplinary and transdisciplinary work. The fifth edition (ECAF'26) will provide a dedicated venue for interdisciplinary discourse on algorithmic fairness via keynotes, paper presentations, a panel discussion, and interactive sessions.

Reference Format:

Tijl De Bie, MaryBeth Defrance, Toon Calders, and Dennis Nguyen . 2026. Fifth European Conference on Algorithmic Fairness (ECAF'26). In *Proceedings of Fifth European Conference on Algorithmic Fairness (ECAF'26)*. Proceedings of Machine Learning Research, 7 pages.

1 Introduction

The Fifth European Conference on Algorithmic Fairness (ECAF'26), formerly known as the European Workshop on Algorithmic Fairness (EWAF), aims to foster dialogue between researchers working on algorithmic fairness in the context of Europe's legal and societal framework, especially in light of the European Union's attempts to promote ethical AI and the turn to AI for the common good. Building upon the success of the first four editions, ECAF'26 will be held from September 2 to September 4, 2026, at Ghent University (Belgium).

2 Objectives and Areas of Interest

The goals of ECAF'26 are to provide a space for algorithmic fairness research grounded in the European perspective and setting, fostering interdisciplinary scholarship and dialogue, and to provide a stimulating scientific program. ECAF welcomes submissions from multiple disciplines, including but not limited to computer science, law, philosophy, and social science. The conference also explicitly welcomes submissions of interdisciplinary and transdisciplinary work, sourcing from multiple disciplinary areas and/or highlighting joint insight-building with relevant non-academic stakeholders. A non-exhaustive list of themes includes:

Authors' Contact Information: Tijl De Bie, Ghent University, Ghent, Belgium, Tijl.Debie@UGent.be; MaryBeth Defrance, Ghent University, Ghent, Belgium, Marybeth.Defrance@UGent.be; Toon Calders, University of Antwerp, Antwerp, Belgium, toon.calders@uantwerpen.be; Dennis Nguyen, Utrecht University, Utrecht, Netherlands, d.nguyen1@uu.nl.

This paper is published under the Creative Commons Attribution-NonCommercial-NoDerivs 4.0 International (CC-BY-NC-ND 4.0) license. Authors reserve their rights to disseminate the work on their personal and corporate Web sites with the appropriate attribution.

ECAF'26, September 02–September 04, 2026, Ghent, BE

© 2026 Copyright held by the owner/author(s).

- Industry experiences in developing and implementing fairness interventions, developing standards and practical approaches to introducing fairness in digital innovation governance.
- Empirical and theoretical perspectives from social sciences on fairness and discrimination in Europe (e.g., analysis of labor markets, the concepts of class, race, disability, and discrimination against minorities in different social contexts, intersectional inequality).
- Case studies based on concrete European instances of algorithmic design and regulation that machine learning scholars or practitioners have encountered in their work (e.g., datasets or audits of automated decision-making systems that are used in Europe).
- An analysis of the implications of the European legislative framework for the debate on fairness in machine learning and AI more broadly (e.g., specificities connected to anti-discrimination and data collection legislation and the emerging regulatory frameworks for platforms and AI).
- Principled arguments for certain fairness concepts and measures in specific contexts.
- Implementing fairness in deployed systems, selecting fairness definitions and designing auditing processes.
- Explorations of the relationship and trade-offs between fairness and transparency in practice.
- Fairness and transparency of black-box models.
- Generative AI and fairness, esp. relating to the job market and the data supply chain.
- Studies on the systemic risks of large online platforms (such as social media platforms and online market-places) and search engines, in particular regarding protection of minors, dissemination of illegal products, gender-based violence, freedom of expression and disinformation.

3 Program

3.1 Accepted Papers

ECAF'26 invited both full papers and extended abstracts. This year, we received 84 submissions, of which 54 extended abstracts and 30 full papers. In total, 69 papers were accepted, of which 48 extended abstracts and 21 full papers. All accepted papers will be presented during either a lightning round talk, poster session, or in-depth presentation.

3.2 Interactive Sessions

Recognizing the crucial role of diverse perspectives in discussions on algorithmic fairness, ECAF seeks contributions that extend beyond the academic discourse, involving voices from industry, civil society, government, NGOs, journalists, artists and other non-academic sectors. To this end, we invited interactive sessions (e.g. workshops, panels, unconferences, moderated dialogues, discussion groups, fish bowls, artistic interventions) to address gaps between theory and practice, research and policy, and multiple disciplines/fields tackling ethical AI; position European institutions, both research and governmental, and their impact on ethical AI; foster conversations between the different stakeholders involved in ethical AI, promoting interdisciplinary and cross-practice discussions; and provide space for discussion and reflection within the algorithmic fairness research community.

This year's edition includes the following accepted interactive sessions:

- "Thinking about AI and humanity @ UGent" by Tijl De Bie

- “Algorithm auditing in practice” by Bram Rijsbosch, Lucas G. Uberti-Bona Marin, Fré Vink, Pauline Remizowa
- “Who Gets the Cake?” An Interactive Exploration of Algorithmic Fairness” by Nina van Liebergen, Eliza Hobo
- “Research and the Digital Services Act: When, why and how to access online platform data” by European Centre for Algorithmic Transparency
- “Verifemmication in action: An intersectional feminist framework for evaluating vulnerability in AI systems” by Anastasia Karagianni
- “Operationalising Fairness and Ethics: Requirements for Governance Engineering in Responsible AI Systems” by Süreyya Akyüz, Dilara Karaduman
- “Industry Panel” by European Association for Algorithmic Fairness
- “Before We Audit Systems, We Audit Ourselves”: A Collective Manifesto for Accountable AI Researchers” by Monique Munarini, Nimra Ahmed, Delaram Golpayegani, Nana Nwachukwu, Miriam Doh
- “What Unfair Treatment Would You Accept from an Algorithmic Manager? Humans’ vs AI as Proposers in the Ultimatum Game” by Jan Grenzebach, Martin Brenzke, Thea Radüntz
- “Beyond Risk: Accountability, Liability, and the Governance of Discriminatory AI in the Global Perspective” by Marilisa D’Amico, Ernesto Damiani, Paolo Ceravolo, Costanza Nardocci
- “Won’t somebody please think of the children? Rebuilding Nuance From Provocations for Children in the Digital Ecosystem” by Hrishita Chakrabarti, Robin Ungruh, Diletta Tobia Micol, Monica Landoni, Sole Pera

3.3 Keynotes

3.3.1 “Is AI Anti-Democratic?” by Mark Coeckelbergh

Abstract. The rapid development and deployment of artificial intelligence presents a profound challenge to the foundational principles of liberal democracy. Far from being politically neutral tools, current AI trajectories are deeply entangled with corporate concentration, surveillance capitalism, and emerging forms of political authoritarianism. Drawing upon recent work on technofascism and the central arguments detailed in *Why AI Undermines Democracy and What to Do About It*, this keynote warns for a structural slide towards technocracy and fascism and examines how AI-driven systems as they are currently developed and used threaten core democratic values such as freedom, fraternity, and tolerance and undermine the epistemic conditions for democracy and civic trust. However, technological fatalism and determinism is rejected. The keynote explores pathways toward democratic resilience. It argues for the democratization of technology itself, the reinvigoration of political institutions, and a renewal of education and civic culture to ensure that digital innovation serves the common good rather than undermining it.

Biography. Mark Coeckelbergh is a leading international philosopher of technology and expert on the ethics of artificial intelligence and robotics. He is Professor of Philosophy of Media and Technology at the University of Vienna and serves as Circle U Chair for Artificial Intelligence. He also holds the ERA Chair at the Institute of Philosophy of the Czech Academy of Sciences, where he leads the Center of Environmental and Technology Ethics. He is a member of the United Nation’s Independent International Scientific Panel on AI, the Belgian Federal Advisory Committee on Data and AI, and previously the High-Level Expert Group on AI of the European

Commission. He is the author of *AI Ethics*, *Introduction to Philosophy of Technology*, *The Political Philosophy of AI*, and *Why AI Undermines Democracy and What To Do About It*.

3.3.2 “*On the Tensions and Trade-offs of Algorithmic Fairness: From Metric Incompatibility to Cross-Pillar Trade-offs*” by Eirini Ntoutsis

Abstract. Algorithmic fairness has evolved into a rich research field. Yet, despite substantial progress, many of its fundamental tensions and trade-offs remain insufficiently understood. These challenges arise at multiple levels. At the level of fairness definitions themselves, impossibility results show that widely used fairness criteria cannot generally be satisfied simultaneously except under restrictive conditions. Additional tensions emerge between individual and group notions of fairness, while intersectional settings reveal a lack of compositionality: ensuring fairness with respect to each protected attribute separately does not guarantee fairness for their combinations. Beyond fairness-specific challenges, fairness must coexist with other core objectives of learning systems. This creates cross-pillar trade-offs between fairness and robustness to data quality issues, adversarial robustness, privacy, and utility. In dynamic environments, these challenges are further compounded by temporal effects, feedback loops, and evolving data distributions that static fairness formulations fail to capture. In this talk, I argue that many of these tensions are currently studied in isolation, obscuring their interactions and cumulative effects. Understanding fairness as one component of a broader ecosystem of often competing objectives is essential for developing responsible AI systems. I conclude by outlining open research challenges and advocating for a more unified research agenda towards responsible AI.

Biography. Eirini Ntoutsis holds the Chair for Open Source Intelligence at the University of the Bundeswehr Munich (UniBwM) and the Research Institute CODE, where she leads the AIML Lab. Her research focuses on adaptive, robust, and responsible AI, including fairness-aware machine learning, bias detection and mitigation, robustness, and learning under dynamic and imperfect data conditions. Her work aims to advance responsible AI systems that deliver positive real-world impact. Prior to joining UniBwM, she was a Professor of Artificial Intelligence at the Free University of Berlin and an Associate Professor of Intelligent Systems at Leibniz University Hannover. She is also affiliated with the L3S Research Center.

3.3.3 “*Limits of Fairness: From Metrics to Practice*” by Mirko Schaefer

Abstract. Established fairness metrics in computer science and responsible AI such as equalized odds, predictive parity, statistical parity, and others, genuinely conflict with one another: optimizing for one systematically costs performance on another. No configuration of these metrics fully secures fairness on its own, and this tension is not a technical shortcoming waiting to be fixed, but a structural one. Additionally, quantification of fairness carries a persuasive promise: certifying fairness with a metric feeds the illusion that fairness checks can be automated, and that quantitative testing can substitute for the qualitative checks and balances that can mitigate what metrics overlook. Fairness cannot be reduced to a metric; it has to be sustained as an ongoing practice. Quantitative bias and model checks are one necessary part, but the context of use, transparency toward stakeholders, and the governance of the system lie beyond what any fairness metric can capture.

This lecture begins by revisiting the limits of fairness metrics, then turns to practices that go beyond quantification. These are often neglected but essential practices that require technical competence and an understanding of ethics, public values and governance. The lecture presents examples for such practices used for responsible procurement,

development, deployment, and governance of algorithmic systems. Drawing on cases from the Netherlands, this lecture shows how these practices increase attention to fairness, constitute and support checks and balances for algorithmic systems. Often inspired or driven by incidents where the use of algorithmic systems has violated social norms and public values, organisations in the Netherlands have developed competences and capacities for responsibly developing, deploying and governing algorithmic systems. The lecture does not dismiss fairness metrics but situates them in a broader understanding of responsible AI. As the Data School has been involved in shaping such practices jointly with practitioners and policy makers, this lecture also demonstrates how collaborative research and stakeholder engagement allows for knowledge transfer, intervention and societal impact.

Biography. Mirko Tobias Schäfer is an Associate Professor of AI, Data & Society at Utrecht University. He leads the interdisciplinary master Applied Data Science at the Department of Information and Computing Sciences. He is co-founder and Science Lead of the Data School at Utrecht University. Together with the Data School, Mirko investigates the social impact of data practices and AI on citizenship and democracy, developing applied methodologies and interventions that bridge research and practice. Mirko also serves on the Advisory Committee Analytics of the Netherlands Ministry of Finance. In addition to his current roles, Mirko was a Visiting Professor at the Helsinki Institute for Social Sciences and Humanities at the University of Helsinki from 2022 to 2025, a Mercator Research Fellow at the NRW School of Governance at the University of Duisburg-Essen in 2016, and a Postdoctoral Researcher at the Centre for the Humanities at Utrecht University from 2014 to 2015.

4 Organization

ECAF'26 was organized by the following people.

General Chairs

- Tijl De Bie, Ghent University
- MaryBeth DeFrance, Ghent University

Program Chairs

- Toon Calders, University of Antwerp
- Dennis Nguyen, Utrecht University

Interactive Sessions Chair

- MaryBeth DeFrance, Ghent University

Web Chair

- Mattia Cerrato, Johannes Gutenberg-Universität Mainz

Proceedings Chairs

- Mattia Cerrato, Johannes Gutenberg-Universität Mainz
- Marco Favier, University of Antwerp

Area Chairs

- Indre Zliobaite, University of Helsinki

- Alessandro Fabris, University of Trieste
- Karima Makhoul, University of Doha for Science and Technology
- Christoph Kern, Ludwig Maximilian University of Munich
- Francien Dechesne, Tilburg University
- Sofie Goethals, University of Antwerp
- Hilde Weerts, Eindhoven University of Technology
- Akra Saxena, Leiden University
- Ludwig Bothmann, Ludwig Maximilian University of Munich
- Eirini Ntoutsis, University of the Bundeswehr Munich

Program Committee

- Andres Algaba, Vrije Universiteit Brussel
- Alessandro Antonucci, IDSIA
- Rhied Al-Othmani, Utrecht University
- Guillaume Bied, Universiteit Gent
- Toon Calders, University of Antwerp
- Roberta Calegari, University of Bologna
- Roger Campdepados, University of Girona
- Marina Ceccon, University of Padua
- Mattia Cerrato, University of Mainz
- Giandomenico Cornacchia, Amazon
- Davide Dalle Pezze, University of Padua
- Marie Decker, Rheinisch Westfälische Technische Hochschule Aachen
- MaryBeth Defrance, Ghent University
- Chadha Degachi, Delft University of Technology
- Eoin D. Delaney, University of Dublin, Trinity College
- Iris Dominguez-Catena, Universidad Pública de Navarra
- Simon Egbert, Universität Bielefeld
- Irmak Erdogan-Peter, KU Leuven
- Marco Favier, Universiteit Antwerpen
- Jan Grenzebach, Federal Institute for Occupational Health and Safety
- Gabriel Grill, Interdisciplinary Transformation University
- Christoph Heitz, ZHAW - Zürcher Hochschule für Angewandte Wissenschaften
- Erik Hekman, Utrecht University
- Corinna Hertweck, University of Zurich
- Ana-Maria Hriscu, Radboud University
- Deborah D. Kanubala, Universität des Saarlandes
- Marius Köppel, ETHZ - ETH Zurich
- Kathrin Lammers, Universität Bielefeld
- Michele Loi, University of Milan

- Blanca Luque Capellas, Johannes-Gutenberg Universität Mainz
- Masoud Mansoury, Delft University of Technology
- Marta Marchiori Manerba, University of Turin
- Fatemeh Mohammadi, University of Milan
- Francesco Nappo, Polytechnic Institute of Milan
- Solal Nathan, Université Paris-Saclay
- Eike Petersen, Fraunhofer Institute for Digital Medicine MEVIS
- Nikolaus Poechhacker, Karl-Franzens-Universität Graz
- Rainer Rehak, Weizenbaum-Institut e.V.
- Emma Rengers, Radboud University
- Theo Röhle, Göteborg University
- Marco Rondina, Polytechnic Institute of Turin
- Raphael Romero, Universiteit Gent
- Benoît Ronval, Ecole Polytechnique de Louvain
- Clara Rus, University of Amsterdam, University of Amsterdam
- Teresa Scantamburlo, University of Venice
- Kristen Marie Scott, KU Leuven
- Ana-Andreea Stoica, Max Planck Institute for Intelligent Systems, Max-Planck Institute
- Chiara Ullstein, Technical University of Munich
- Flore Vancompernelle Vromman, LouRIM
- Madeleine Waller, Eberhard-Karls-Universität Tübingen
- Nick Wils, University of Antwerp
- Farjana Yesmin, Independent Researcher
- Sebastian Zezulka, Eberhard-Karls-Universität Tübingen

Acknowledgments

ECAF'26 was made possible by the generous support of the European Centre for Algorithmic Transparency (ECAT), the Flemish Government, and the Doctoral School of Ghent University.