

A EU Icon Categories and Regulatory Thresholds

Table 1. EU icon categories in relation to the AI Act, the DSA, the Code of Practice, and the threshold analysis developed in this paper.

Category	AI Act	DSA	Code of Practice	Conceptual Dimension
Basic AI icon	Articles 50(4) and 50(5) require deployers to disclose qualifying deepfakes and public-interest text clearly, distinguishably, accessibly, and no later than first exposure.	Article 35(1)(k) refers to prominent interface markings and easy-to-use functionality through which recipients can indicate generated or manipulated content.	The basic icon may accompany a customised textual disclosure or an interactive second layer. The Code specifies design and placement requirements for such disclosures [1, pp. 29–33, 37].	Supports context-specific disclosure where a generic label would not sufficiently explain which element was generated or modified.
Fully generated “AI GENERATED”	Article 50(2) requires provider-side machine-readable marking and detectability. Article 50(4) adds visible disclosure where the output constitutes a deepfake or qualifying public-interest text.	Article 35(1)(k) applies where generated image, audio, or video appreciably resembles existing persons, objects, places, entities, or events and falsely appears authentic or truthful.	The Code provides for digitally signed metadata, imperceptible watermarking, and detection mechanisms. The Annex identifies the “AI GENERATED” icon for fully generated content [1, pp. 8–15, 37].	Represents the clearest case in which content is “AI-generated enough.” Remaining questions concern resemblance, false authenticity, publication purpose, and statutory exceptions.
Partially modified “AI MODIFIED”	Articles 3(60) and 50(4) cover AI-generated or manipulated media that resembles existing persons, objects, places, entities, or events and falsely appears authentic or truthful.	Article 35(1)(k) similarly covers materially generated or manipulated media satisfying the resemblance and false-authenticity criteria.	The “AI MODIFIED” icon applies to pre-existing human-created content that has been partially modified with AI. A second layer may explain the nature and location of the modification [1, pp. 33, 37].	Central to the threshold problem. Disclosure should depend on whether AI materially changes a represented person, event, statement, or evidentiary claim.
AI-assisted or minor enhancement	Article 50 does not make every use of AI subject to visible disclosure. Article 50(4) depends on whether the content constitutes a deepfake or qualifying manipulated public-interest text.	Article 35(1)(k) is not triggered merely because AI was used; the resemblance and false-authenticity requirements must also be satisfied.	The Code does not establish a separate icon for minor correction, noise reduction, automated subtitles, or limited enhancement, but encourages richer provenance and layered information [1, pp. 11, 33, 37].	Illustrates the unresolved boundary. Challenges with how disclosure should depend on materiality, communicative significance, and the risk of false authenticity rather than mere AI involvement.
Recipient indication	The primary duties remain with providers and deployers; recipients are not responsible for determining whether content legally qualifies as a deepfake.	Article 35(1)(k) requires easy-to-use recipient functionality. Articles 16 and 22 may also be relevant where notices concern illegal content.	Detection solutions should be available to users, researchers, civil society, media, fact-checkers, trusted flaggers, and authorities [1, pp. 12–15].	Creates interface and responsibility-design problems. Users need usable categories, explanations, detection support, and mechanisms to challenge erroneous decisions.
Machine-readable provenance	Article 50(2) requires machine-readable marking and detectability, separately from visible deployer disclosure under Article 50(4).	Technical provenance may support Article 35 mitigation, but invisible metadata alone does not necessarily satisfy the objective of prominent interface disclosure.	The Code provides for digitally signed metadata, imperceptible watermarking, optional fingerprinting or logging, and richer provenance information [1, pp. 8–11].	Shows that provenance and visible labelling perform different functions. Technical origin information does not by itself determine whether content is misleading, harmful, satirical, or legally subject to disclosure.

B Statement of Reasons Dashboard Screenshots




Fig. 1. Filter settings used for the Statement of Reasons dashboard analysis. The dashboard is filtered to the content type “Synthetic Media” for the period from 25 September 2023 to 28 April 2026, with all platforms, restrictions, information sources, violations, grounds for restriction, decision automation, and automated detection categories included.




Fig. 2. Statements of Reasons for “Synthetic Media” by platform. The filtered view shows 153,138,573 statements of reasons in total, with most entries concentrated on Facebook, followed by Instagram, X, and Threads.




Fig. 3. Restrictions and violation categories for “Synthetic Media” statements of reasons. The left chart shows that visibility restrictions dominate the reported enforcement measures. The right treemap shows the distribution of reported violation categories, including data protection and privacy violations, scope of platform service, illegal or harmful speech, violence, pornography or sexualized content, and scams or fraud.

Fig. 4. Daily submissions of “Synthetic Media” Statements of Reasons by violation category. The upper chart displays daily submission volumes, while the lower chart shows changes in the proportional distribution of reported violation categories. The temporal axis is reversed in the dashboard. The dashboard snapshot was collected on 30 April 2026.




Fig. 5. Daily submissions of “Synthetic Media” statements of reasons by platform. The temporal axis is reversed in the dashboard. The figure shows a pronounced spike in Facebook submissions around late April 2026, while other platforms remain comparatively low across the displayed period.




Fig. 6. Daily submissions of “Synthetic Media” statements of reasons by violation category. The dashboard shows changes over time in reported violation categories, with a visible spike in April 2026 and a changing distribution of violation categories across the displayed days.




Fig. 7. Statements of Reasons for “Synthetic Media” by restriction type. Removal of content is the dominant reported restriction, followed by demotion of content. Labelled content appears as a much smaller category, indicating that synthetic-media moderation is reported more often through removal or demotion than through labelling.